

Quiz2 Machine Learning

Fatemeh Abyarjoo

Question 1:

```
# install.packages("AppliedPredictiveModeling")
# install.packages("ggplot2")
# install.packages("lattice")
library(lattice)
library(ggplot2)
library(AppliedPredictiveModeling)
library(caret)
data(AlzheimerDisease)

# adData = data.frame(diagnosis,predictors)
# testIndex = createDataPartition(diagnosis, p = 0.50,list=FALSE)
# training = adData[-testIndex,]
# testing = adData[testIndex,]
```

Question 2: Load the cement data. Make a histogram and confirm the SuperPlasticizer variable is skewed. Normally you might use the log transform to try to make the data more symmetric. Why would that be a poor choice for this variable?

```
library(AppliedPredictiveModeling)
data(concrete)
library(caret)
set.seed(1000)
inTrain = createDataPartition(mixtures$CompressiveStrength, p = 3/4)[[1]]
training = mixtures[ inTrain,]
testing = mixtures[-inTrain,]

names(training)
```


```
## [1] "Cement"          "BlastFurnaceSlag"  "FlyAsh"
## [4] "Water"           "Superplasticizer"  "CoarseAggregate"
## [7] "FineAggregate"   "Age"               "CompressiveStrength"
```

```
summary(training)
```

```
##      Cement      BlastFurnaceSlag      FlyAsh      Water
## Min.   :0.04482 Min.   :0.000000 Min.   :0.00000 Min.   :0.05139
## 1st Qu.:0.08179 1st Qu.:0.000000 1st Qu.:0.00000 1st Qu.:0.06972
## Median :0.11462 Median :0.009993 Median :0.00000 Median :0.07862
## Mean   :0.11782 Mean   :0.032051 Mean   :0.02247 Mean   :0.07774
## 3rd Qu.:0.14793 3rd Qu.:0.061968 3rd Qu.:0.04999 3rd Qu.:0.08384
## Max.   :0.22541 Max.   :0.150339 Max.   :0.08884 Max.   :0.11222
## Superplasticizer CoarseAggregate FineAggregate Age
## Min.   :0.000000 Min.   :0.3459 Min.   :0.2480 Min.   : 1.00
## 1st Qu.:0.000000 1st Qu.:0.3986 1st Qu.:0.3113 1st Qu.: 14.00
## Median :0.002726 Median :0.4213 Median :0.3305 Median : 28.00
```


```
## Mean :0.002608 Mean :0.4167 Mean :0.3306 Mean : 47.46
## 3rd Qu.:0.004351 3rd Qu.:0.4389 3rd Qu.:0.3542 3rd Qu.: 56.00
## Max. :0.013149 Max. :0.4798 Max. :0.4141 Max. :365.00
## CompressiveStrength
## Min. : 2.33
## 1st Qu.:23.71
## Median :34.48
## Mean :35.64
## 3rd Qu.:46.13
## Max. :82.60
```

```
# install.packages("ISLR")
library(ISLR)
featurePlot(x = training[,c("Cement", "BlastFurnaceSlag", "FlyAsh", "Water",
                             "Superplasticizer", "CoarseAggregate", "FineAggregate",
                             "Age", "CompressiveStrength")], y = training$CompressiveStrength,
            plot = "pairs")
```

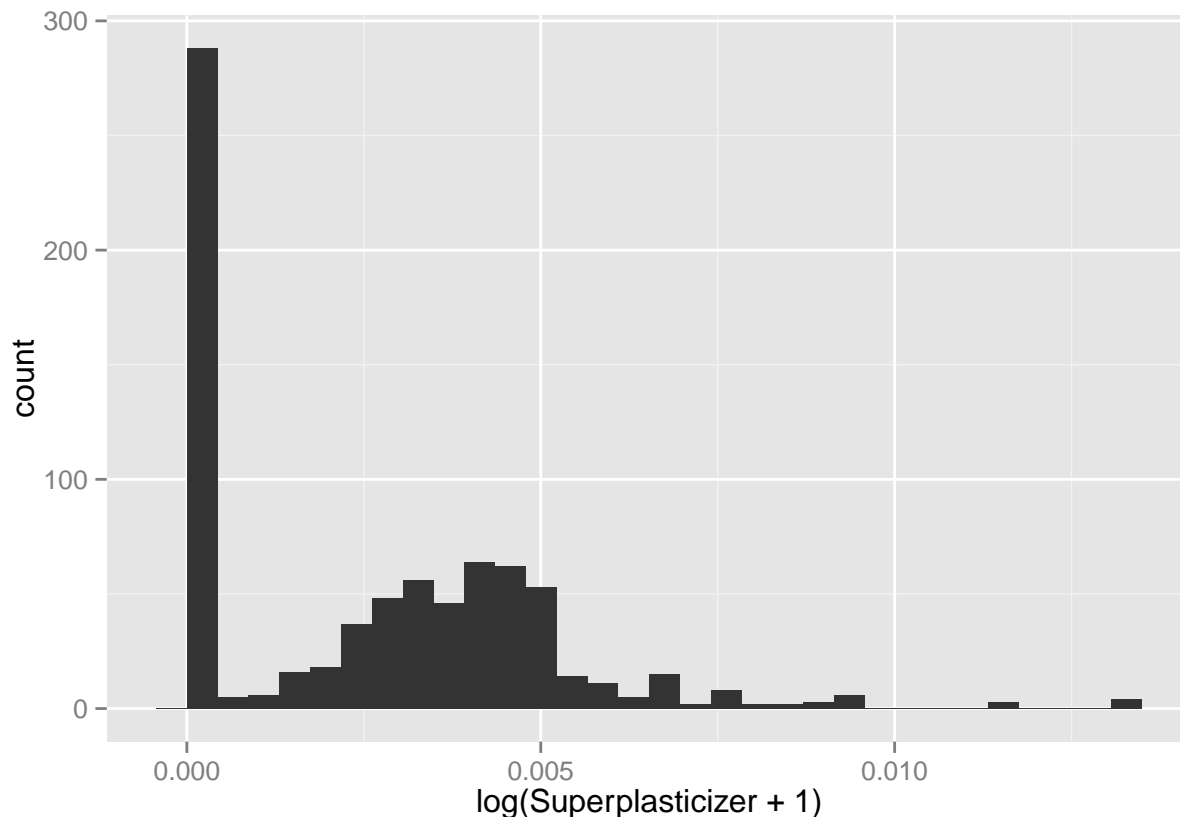


Scatter Plot Matrix

```
hist(training$Superplasticizer)
```



```
qplot(log(Superplasticizer+1),data=training)
```



There are a large number of values that are the same and even if you took the $\log(\text{SuperPlasticizer} + 1)$ they would still all be identical so the distribution would not be symmetric.

Question 3 Load the Alzheimer's disease data. Find all the predictor variables in the training set that begin with IL. Perform principal components on these variables with the `preProcess()` function from the `caret` package. Calculate the number of principal components needed to capture 90% of the variance. How many are there?

```
library(caret)
library(AppliedPredictiveModeling)
set.seed(3433)
data(AlzheimerDisease)
adData = data.frame(diagnosis, predictors)
inTrain = createDataPartition(adData$diagnosis, p = 3/4)[[1]]
training = adData[ inTrain,]
testing = adData[-inTrain,]

names(training)
```

```
## [1] "diagnosis"
## [3] "ACTH_Adrenocorticotrophic_Hormon"
## [5] "Adiponectin"
## [7] "Alpha_1_Antitrypsin"
## [9] "Alpha_2_Macroglobulin"
## [11] "Angiotensinogen"
## [13] "Apolipoprotein_A1"
## [15] "Apolipoprotein_B"
## [17] "ACE_CD143_Angiotensin_Converti"
## [19] "AXL"
## [21] "Alpha_1_Antichymotrypsin"
## [23] "Alpha_1_Microglobulin"
## [25] "Angiopoietin_2_ANG_2"
## [27] "Apolipoprotein_A_IV"
## [29] "Apolipoprotein_A2"
## [31] "Apolipoprotein_CI"
```

| | | |
|----------|-----------------------------------|------------------------------------|
| ## [17] | "Apolipoprotein_CIII" | "Apolipoprotein_D" |
| ## [19] | "Apolipoprotein_E" | "Apolipoprotein_H" |
| ## [21] | "B_Lymphocyte_Chemoattractant_BL" | "BMP_6" |
| ## [23] | "Beta_2_Microglobulin" | "Betacellulin" |
| ## [25] | "C_Reactive_Protein" | "CD40" |
| ## [27] | "CD5L" | "Calbindin" |
| ## [29] | "Calcitonin" | "CgA" |
| ## [31] | "Clusterin_Apo_J" | "Complement_3" |
| ## [33] | "Complement_Factor_H" | "Connective_Tissue_Growth_Factor" |
| ## [35] | "Cortisol" | "Creatine_Kinase_MB" |
| ## [37] | "Cystatin_C" | "EGF_R" |
| ## [39] | "EN_RAGE" | "ENA_78" |
| ## [41] | "Eotaxin_3" | "FAS" |
| ## [43] | "FSH_Follicle_Stimulation_Hormon" | "Fas_Ligand" |
| ## [45] | "Fatty_Acid_Binding_Protein" | "Ferritin" |
| ## [47] | "Fetuin_A" | "Fibrinogen" |
| ## [49] | "GRO_alpha" | "Gamma_Interferon_induced_Monokin" |
| ## [51] | "Glutathione_S_Transferase_alpha" | "HB_EGF" |
| ## [53] | "HCC_4" | "Hepatocyte_Growth_Factor_HGF" |
| ## [55] | "I_309" | "ICAM_1" |
| ## [57] | "IGF_BP_2" | "IL_11" |
| ## [59] | "IL_13" | "IL_16" |
| ## [61] | "IL_17E" | "IL_1alpha" |
| ## [63] | "IL_3" | "IL_4" |
| ## [65] | "IL_5" | "IL_6" |
| ## [67] | "IL_6_Receptor" | "IL_7" |
| ## [69] | "IL_8" | "IP_10_Inducible_Protein_10" |
| ## [71] | "IgA" | "Insulin" |
| ## [73] | "Kidney_Injury_Molecule_1_KIM_1" | "LOX_1" |
| ## [75] | "Leptin" | "Lipoprotein_a" |
| ## [77] | "MCP_1" | "MCP_2" |
| ## [79] | "MIF" | "MIP_1alpha" |
| ## [81] | "MIP_1beta" | "MMP_2" |
| ## [83] | "MMP_3" | "MMP10" |
| ## [85] | "MMP7" | "Myoglobin" |
| ## [87] | "NT_proBNP" | "NrCAM" |
| ## [89] | "Osteopontin" | "PAI_1" |
| ## [91] | "PAPP_A" | "PLGF" |
| ## [93] | "PYY" | "Pancreatic_polypeptide" |
| ## [95] | "Prolactin" | "Prostatic_Acid_Phosphatase" |
| ## [97] | "Protein_S" | "Pulmonary_and_Activation_Regulat" |
| ## [99] | "RANTES" | "Resistin" |
| ## [101] | "S100b" | "SGOT" |
| ## [103] | "SHBG" | "SOD" |
| ## [105] | "Serum_Amyloid_P" | "Sortilin" |
| ## [107] | "Stem_Cell_Factor" | "TGF_alpha" |
| ## [109] | "TIMP_1" | "TNF_RII" |
| ## [111] | "TRAIL_R3" | "TTR_prealbumin" |
| ## [113] | "Tamm_Horsfall_Protein_THP" | "Thrombomodulin" |
| ## [115] | "Thrombopoietin" | "Thymus_Expressed_Chemokine_TECK" |
| ## [117] | "Thyroid_Stimulating_Hormone" | "Thyroxine_Binding_Globulin" |
| ## [119] | "Tissue_Factor" | "Transferrin" |
| ## [121] | "Trefoil_Factor_3_TFF3" | "VCAM_1" |
| ## [123] | "VEGF" | "Vitronectin" |

```
## [125] "von_Willebrand_Factor"      "age"
## [127] "tau"                        "p_tau"
## [129] "Ab_42"                     "male"
## [131] "Genotype"
```

```
ILset=grep("^IL", names(training), value = TRUE)

Nofcom =preProcess(training[, ILset], method = "pca", thresh = 0.9)
Nofcom
```

```
##
## Call:
## preProcess.default(x = training[, ILset], method = "pca", thresh = 0.9)
##
## Created from 251 samples and 12 variables
## Pre-processing: principal component signal extraction, scaled, centered
##
## PCA needed 9 components to capture 90 percent of the variance
```

Question 4: Load the Alzheimer's disease data. Create a training data set consisting of only the predictors with variable names beginning with IL and the diagnosis. Build two predictive models, one using the predictors as they are and one using PCA with principal components explaining 80% of the variance in the predictors. Use method="glm" in the train function. What is the accuracy of each method in the test set? Which is more accurate?

```
library(lattice)
library(ggplot2)
library(caret)
library(AppliedPredictiveModeling)
set.seed(3433)
data(AlzheimerDisease)
adData = data.frame(diagnosis,predictors)
inTrain = createDataPartition(adData$diagnosis, p = 3/4)[[1]]
training = adData[ inTrain,]
testing = adData[-inTrain,]

set.seed(3433)
ILset=grep("^IL", names(training), value = TRUE)
ILpredictor= predictors[, ILset]
dataset= data.frame(diagnosis, ILpredictor)
inTrain = createDataPartition(dataset$diagnosis, p = 3/4)[[1]]
training = dataset[inTrain, ]
testing = dataset[-inTrain, ]

# install.packages("Hmisc")
# install.packages("survival")
# install.packages("gridExtra")
# install.packages("dplyr")
# install.packages('e1071', dependencies=TRUE)

library(gridExtra)
```

```
## Loading required package: grid
```

```
library(survival)
```

```
##
## Attaching package: 'survival'
##
## The following object is masked from 'package:caret':
##
##   cluster
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
##
## The following object is masked from 'package:stats':
##
##   filter
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(Hmisc)
```

```
## Loading required package: Formula
##
## Attaching package: 'Hmisc'
##
## The following objects are masked from 'package:dplyr':
##
##   combine, src, summarize
##
## The following objects are masked from 'package:base':
##
##   format.pval, round.POSIXt, trunc.POSIXt, units
```

```
Mod1=train(diagnosis ~ ., method = "glm", data = training)
predictions=predict(Mod1, newdata = testing)
Confusionmat1= confusionMatrix(predictions, testing$diagnosis)
print(Confusionmat1)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction Impaired Control
##   Impaired      2      9
##   Control     20     51
##
##           Accuracy : 0.6463
##           95% CI : (0.533, 0.7488)
##   No Information Rate : 0.7317
```

```
##      P-Value [Acc > NIR] : 0.96637
##
##              Kappa : -0.0702
## Mcnemar's Test P-Value : 0.06332
##
##      Sensitivity : 0.09091
##      Specificity : 0.85000
##      Pos Pred Value : 0.18182
##      Neg Pred Value : 0.71831
##      Prevalence : 0.26829
##      Detection Rate : 0.02439
##      Detection Prevalence : 0.13415
##      Balanced Accuracy : 0.47045
##
##      'Positive' Class : Impaired
##
```

```
Mod2=train(training$diagnosis ~ ., method = "glm", preProcess = "pca",
  data = training, trControl = trainControl(preProcOptions = list(thresh = 0.8)))
Confusionmat2=confusionMatrix(testing$diagnosis, predict(Mod2, testing))
print(Confusionmat2)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction Impaired Control
##   Impaired      3      19
##   Control       4      56
##
##      Accuracy : 0.7195
##      95% CI : (0.6094, 0.8132)
##      No Information Rate : 0.9146
##      P-Value [Acc > NIR] : 1.000000
##
##      Kappa : 0.0889
## Mcnemar's Test P-Value : 0.003509
##
##      Sensitivity : 0.42857
##      Specificity : 0.74667
##      Pos Pred Value : 0.13636
##      Neg Pred Value : 0.93333
##      Prevalence : 0.08537
##      Detection Rate : 0.03659
##      Detection Prevalence : 0.26829
##      Balanced Accuracy : 0.58762
##
##      'Positive' Class : Impaired
##
```